

<div>on RequestVote request from peer</div> <pre> if currentTerm &lt; m.term:   stepDown(m.term) if (currentTerm == m.term and   votedFor in [None, peer] and   (m.lastLogTerm &gt; logTerm(len(log)) or   (m.lastLogTerm == logTerm(len(log)) and   m.lastLogIndex &gt;= len(log))):   granted := True   votedFor = peer   electionAlarm = now() + rand(1.0, 2.0) *   ELECTION_TIMEOUT else:   granted := False reply {term: currentTerm,   granted: granted} </pre> <div>on RequestVote response from peer</div> <pre> if currentTerm &lt; m.term:   stepDown(m.term) if (state == CANDIDATE and   currentTerm == m.term):   rpcDue[peer] = INFINITY   voteGranted[peer] = m.granted </pre> <div>on AppendEntries request from peer</div> <pre> if currentTerm &lt; m.term:   stepDown(m.term) if currentTerm &gt; m.term:   reply {term: currentTerm,   success: False} else:   leader = peer   state = FOLLOWER   electionAlarm = now() + rand(1.0, 2.0) *   ELECTION_TIMEOUT   success := (m.prevIndex == 0 or   (m.prevIndex &lt;= len(log) and   log[m.prevIndex].term == m.prevTerm))   if success:     index := m.prevIndex     for j := 1..len(m.entries):       index += 1       if getTerm(index) != m.entries[j].term:       log = log[1..(index-1)] + m.entries[j]     commitIndex = min(m.commitIndex, index)   else:     index = 0   reply {term: currentTerm,   success: success,   matchIndex: index} </pre> <div>on AppendEntries response from peer</div> <pre> if currentTerm &lt; m.term:   stepDown(m.term) elif state == LEADER and currentTerm == m.term:   if m.success:     matchIndex[peer] = m.matchIndex     nextIndex[peer] = m.matchIndex + 1   else:     nextIndex[peer] = max(1, nextIndex[peer] - 1) </pre> <div>on StateMachine request from client</div> <pre> if state == LEADER:   log.append({term: currentTerm,   command: m.command}) </pre> <div>helper functions</div> <pre> def stepDown(newTerm):   currentTerm = newTerm   state = FOLLOWER   votedFor = None   if electionAlarm &lt; now():     electionAlarm = now() + rand(1.0, 2.0) *     ELECTION_TIMEOUT def logTerm(index):   if index &lt; 1 or index &gt; len(log):     return 0   else:     return log[index].term </pre>	<div>start new election</div> <pre> on (state in [FOLLOWER, CANDIDATE] and   electionAlarm &lt; now()):   electionAlarm = now() + rand(1.0, 2.0) *   ELECTION_TIMEOUT   currentTerm += 1   votedFor = serverID   state = CANDIDATE   foreach peer:     # reset all state for peer </pre> <div>send RequestVote to peer</div> <pre> on (state == CANDIDATE and   rpcDue[peer] &lt; now()):   rpcDue[peer] = now() + RPC_TIMEOUT   send RequestVote to peer {   term: currentTerm,   lastLogTerm: logTerm(len(log)),   lastLogIndex: len(log)} </pre> <div>become leader</div> <pre> on (state == CANDIDATE and   sum(voteGranted) + 1 &gt; NUM_SERVERS / 2:   state = LEADER   leader = localhost   foreach peer:     nextIndex[peer] = len(log) + 1 </pre> <div>send AppendEntries to peer</div> <pre> on (state == LEADER and   (matchIndex[peer] &lt; len(log) or   rpcDue[peer] &lt; now()):   rpcDue[peer] = now() + ELECTION_TIMEOUT / 2   lastIndex := choose in (nextIndex[peer] - 1)..len(log)   nextIndex[peer] = lastIndex   send AppendEntries to peer {   term: currentTerm,   prevIndex: nextIndex[peer] - 1,   prevTerm: getTerm(nextIndex[peer] - 1),   entries: log[nextIndex[peer]..lastIndex],   commitIndex: commitIndex} </pre> <div>advance commit index</div> <pre> n := sorted(matchIndex + [len(log)])[NUM_SERVERS / 2ish] on (state == LEADER and   logTerm(n) == currentTerm):   commitIndex = n </pre> <div>advance state machine</div> <pre> on lastApplied &lt; commitIndex:   lastApplied += 1   result := stateMachine.apply(log[lastApplied])   if (state == Leader and   logTerm(lastApplied) == currentTerm):     # send result to client </pre> <div>server state</div> <pre> # FOLLOWER, CANDIDATE, or LEADER state := FOLLOWER # latest term server has seen (increases monotonically) currentTerm := 1 # candidate that received vote in current term votedFor := None # log entries; each entry contains command for state # machine, and term when entry was received by leader log := [] # indexed from 1 # index of highest log entry known to be committed commitIndex := 0 # time after which to start new election electionAlarm := 0.0 # applies committed commands in log order stateMachine := new SM() # identity of last known leader leader := None  # State per peer, valid only for the current term foreach peer:   # time after which to send another RPC   # (RequestVote or heartbeat)   rpcDue[peer] := 0.0   # True if peer has granted this server its vote   voteGranted[peer] := False   # index of highest log entry known to be replicated   # on peer   matchIndex[peer] := 0   # index of next log entry to send to peer   nextIndex[peer] := 1 </pre>
--	--

Figure 1: alternative cheatsheet